

▼ Electricity Price Prediction with Machine Learning

The price of electricity depends on many factors. Predicting the price of electricity helps many businesses understand how much electricity they have to pay each year. The Electricity Price Prediction task is based on a case study where you need to predict the daily price of electricity based on the daily consumption of heavy machinery used by businesses.

Suppose that your business relies on computing services where the power consumed by your machines varies throughout the day. You do not know the actual cost of the electricity consumed by the machines throughout the day, but the organization has provided you with historical data of the price of the electricity consumed by the machines. Below is the information of the data we have for the task of forecasting electricity prices:

DateTime: Date and time of the record

Holiday: contains the name of the holiday if the day is a national holiday

HolidayFlag: contains 1 if it's a bank holiday otherwise 0

DayOfWeek: contains values between 0-6 where 0 is Monday

WeekOfYear: week of the year

Day: Day of the date

Month: Month of the date

Year: Year of the date

PeriodOfDay: half-hour period of the day

ForecastWindProduction: forecasted wind production

SystemLoadEA forecasted national load

SMPEA: forecasted price

ORKTemperature: actual temperature measured

ORKWindspeed: actual windspeed measured

CO2Intensity: actual CO2 intensity for the electricity produced

ActualWindProduction: actual wind energy production

SystemLoadEP2: actual national system load

SMPEP2: the actual price of the electricity consumed (labels or values to be predicted)

So my task here is to use this data to train a machine learning model to predict the price of electricity consumed by the machines.

I will start the task of electricity price prediction by importing the necessary Python libraries and the dataset that we need for this task:

```
import pandas as pd
import numpy as np
data = pd.read_csv("Electricity_data.csv")
print(data.head())
```

	DateTime	Holiday	HolidayFlag	DayOfWeek	WeekOfYear	Day	Month	\
0	01/11/2011 00:00	None	0	1	44	1	11	

1	01/11/2011	00:30	None	0	1	44	1	11
2	01/11/2011	01:00	None	0	1	44	1	11
3	01/11/2011	01:30	None	0	1	44	1	11
4	01/11/2011	02:00	None	0	1	44	1	11

	Year	PeriodOfDay	ForecastWindProduction	SystemLoadEA	SMPEA	\
0	2011	0	315.31	3388.77	49.26	
1	2011	1	321.80	3196.66	49.26	
2	2011	2	328.57	3060.71	49.10	
3	2011	3	335.60	2945.56	48.04	
4	2011	4	342.90	2849.34	33.75	

	ORKTemperature	ORKWindspeed	CO2Intensity	ActualWindProduction	SystemLoadEP2	\
0	6.00	9.30	600.71		3159.60	
1	6.00	11.10	605.42		2973.01	
2	5.00	11.10	589.97		2834.00	
3	6.00	9.30	585.94		2725.99	
4	6.00	11.10	571.52		2655.64	

	SMPEP2
0	54.32
1	54.23
2	54.23
3	53.47
4	39.87

```
<ipython-input-1-80d0d62a73e2>:3: DtypeWarning: Columns (9,10,11,14,15,16,17) have mixed types. Specify
data = pd.read_csv("Electricity_data.csv")
```

Let's have a look at all the columns of this dataset:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38014 entries, 0 to 38013
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   DateTime                             38014 non-null  object
1   Holiday                             38014 non-null  object
2   HolidayFlag                          38014 non-null  int64
3   DayOfWeek                           38014 non-null  int64
4   WeekOfYear                           38014 non-null  int64
5   Day                                  38014 non-null  int64
6   Month                                38014 non-null  int64
7   Year                                 38014 non-null  int64
8   PeriodOfDay                          38014 non-null  int64
9   ForecastWindProduction               38014 non-null  object
10  SystemLoadEA                         38014 non-null  object
11  SMPEA                                38014 non-null  object
12  ORKTemperature                       38014 non-null  object
13  ORKWindspeed                         38014 non-null  object
14  CO2Intensity                         38014 non-null  object
15  ActualWindProduction                 38014 non-null  object
16  SystemLoadEP2                       38014 non-null  object
17  SMPEP2                              38014 non-null  object
dtypes: int64(7), object(11)
memory usage: 5.2+ MB
```

I can see that so many features with numerical values are string values in the dataset and not integers or float values. So before moving further, we have to convert these string values to float values:

```
data["ForecastWindProduction"] = pd.to_numeric(data["ForecastWindProduction"], errors= 'coerce')
data["SystemLoadEA"] = pd.to_numeric(data["SystemLoadEA"], errors= 'coerce')
data["SMPEA"] = pd.to_numeric(data["SMPEA"], errors= 'coerce')
data["ORKTemperature"] = pd.to_numeric(data["ORKTemperature"], errors= 'coerce')
```

```
data["ORKWindspeed"] = pd.to_numeric(data["ORKWindspeed"], errors= 'coerce')
data["CO2Intensity"] = pd.to_numeric(data["CO2Intensity"], errors= 'coerce')
data["ActualWindProduction"] = pd.to_numeric(data["ActualWindProduction"], errors= 'coerce')
data["SystemLoadEP2"] = pd.to_numeric(data["SystemLoadEP2"], errors= 'coerce')
data["SMPEP2"] = pd.to_numeric(data["SMPEP2"], errors= 'coerce')
```

Now let's have a look at whether this dataset contains any null values or not:

```
data.isnull().sum()
```

```
DateTime          0
Holiday           0
HolidayFlag       0
DayOfWeek         0
WeekOfYear        0
Day              0
Month            0
Year            0
PeriodOfDay      0
ForecastWindProduction  5
SystemLoadEA      2
SMPEA            2
ORKTemperature    295
ORKWindspeed      299
CO2Intensity       7
ActualWindProduction  5
SystemLoadEP2     2
SMPEP2           2
dtype: int64
```

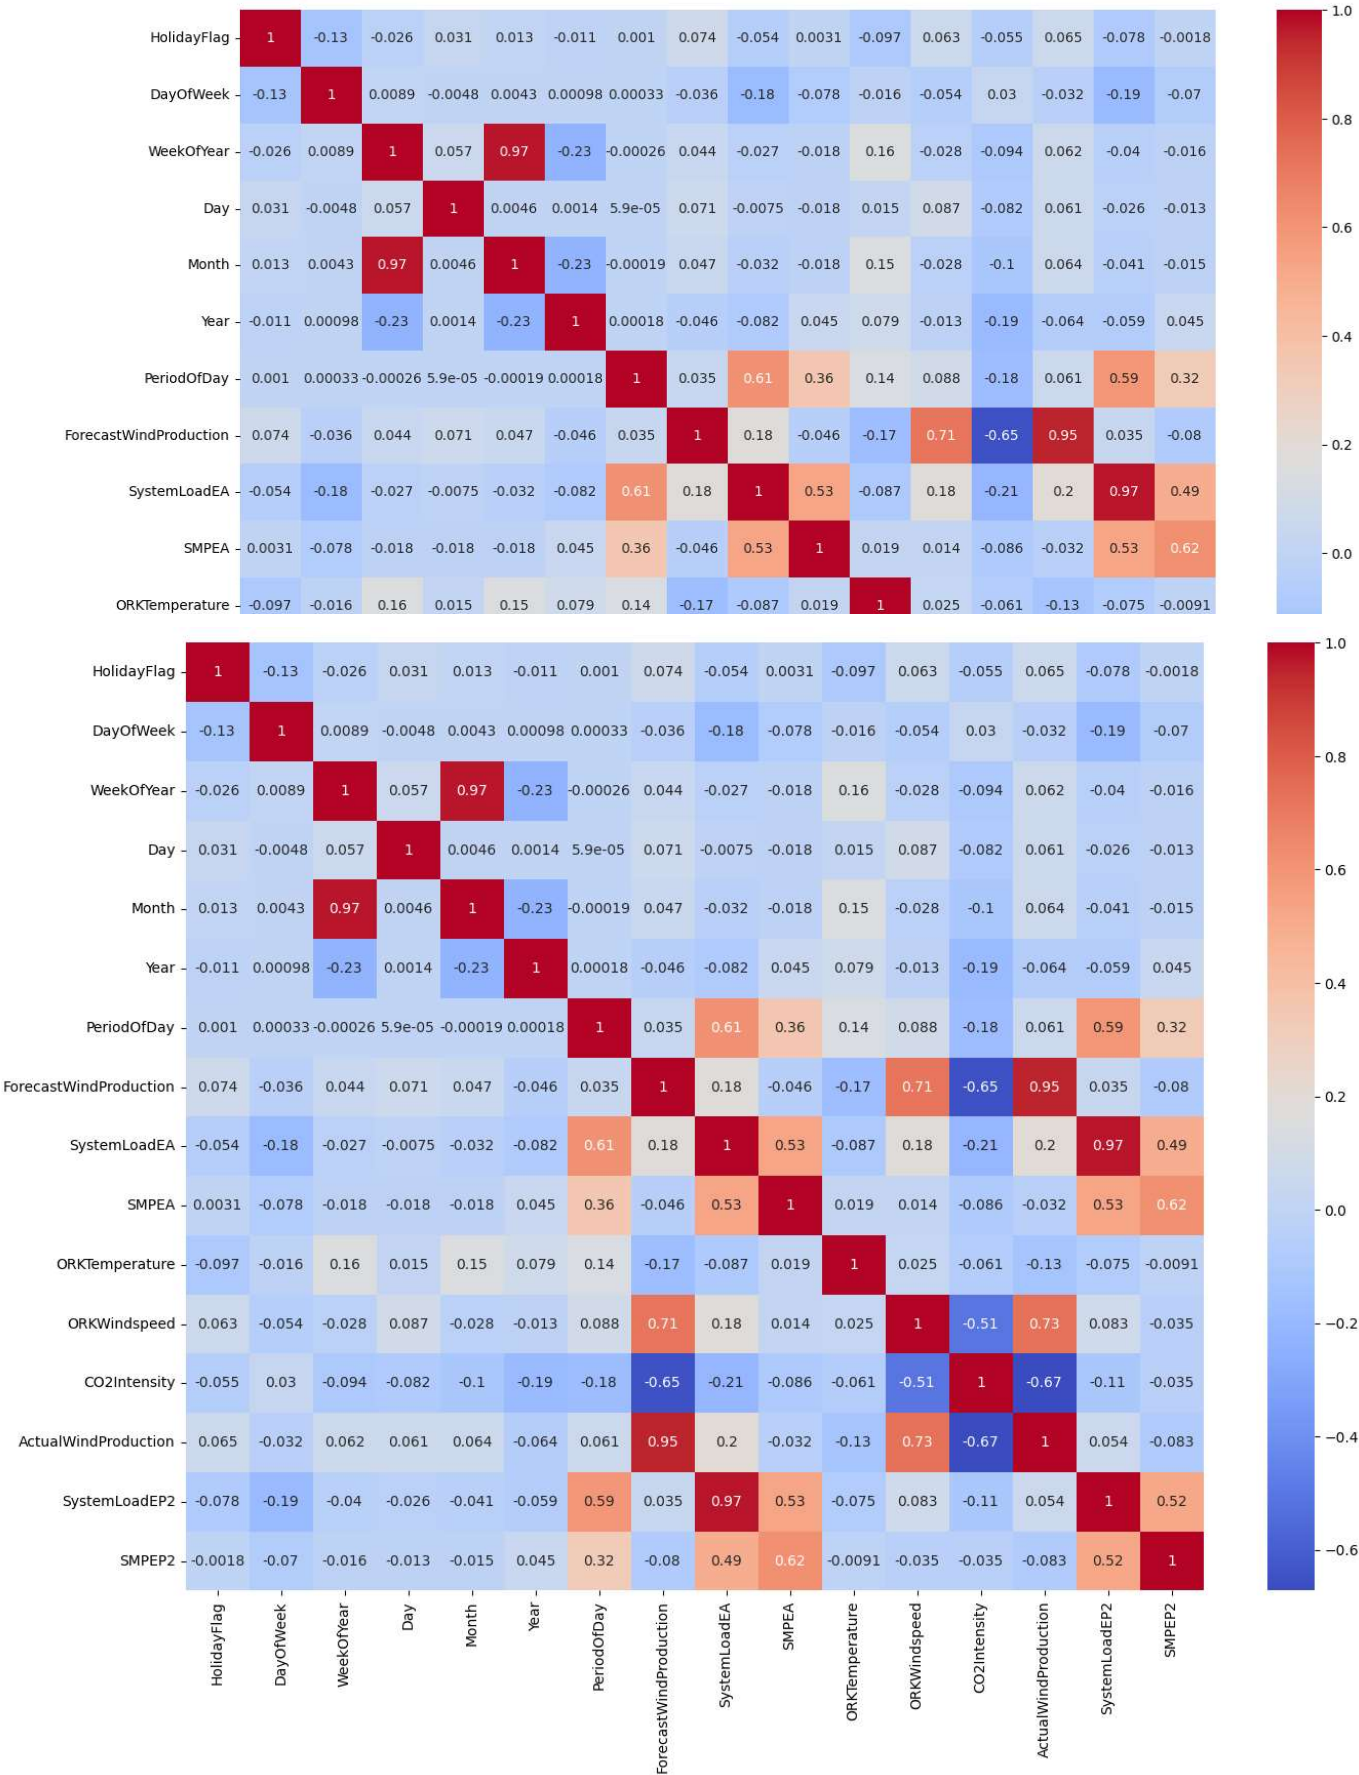
So there are some columns with null values, I will drop all these rows containing null values from the dataset:

```
data = data.dropna()
```

Now let's have a look at the correlation between all the columns in the dataset:

```
import seaborn as sns
import matplotlib.pyplot as plt
correlations = data.corr(method='pearson')
plt.figure(figsize=(16, 12))
sns.heatmap(correlations, cmap="coolwarm", annot=True)
plt.show()
```

```
<ipython-input-6-c51eeb38dbc7>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is correlations = data.corr(method='pearson')
```



Electricity Price Prediction Model

Now let's move to the task of training an electricity price prediction model. Here I will first add all the important features to x and the target column to y, and then I will split the data into training and test sets:

```
x = data[["Day", "Month", "ForecastWindProduction", "SystemLoadEA",
          "SMPEA", "ORKTemperature", "ORKWindspeed", "CO2Intensity",
          "ActualWindProduction", "SystemLoadEP2"]]
y = data["SMPEP2"]
from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(x, y,
                                              test_size=0.2,
                                              random_state=42)
```

As this is a regression problem, so here I will choose the Random Forest regression algorithm to train the electricity price prediction model:

```
from sklearn.ensemble import RandomForestRegressor
model = RandomForestRegressor()
model.fit(xtrain, ytrain)
```

▼ RandomForestRegressor

RandomForestRegressor()

Now let's input all the values of the necessary features that we used to train the model and have a look at the price of the electricity predicted by the model:

```
#features = [ ["Day", "Month", "ForecastWindProduction", "SystemLoadEA", "SMPEA", "ORKTemperature", "ORKWindspeed", "CO2Intensity", "ActualWindProduction", "SystemLoadEP2"] ]
features = np.array([[10, 12, 54.10, 4241.05, 49.56, 9.0, 14.8, 491.32, 54.0, 4426.84]])
model.predict(features)

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, using indices instead
  warnings.warn(
array([67.3126])
```

